QuIT your B$^+$-tree for the *Quick Insertion Tree*

Aneesh Raman
Boston University
aneeshr@bu.edu

Konstantinos Karatsenidis
Boston University
karatse@bu.edu

Shaolin Xie
Boston University
slxie@bu.edu

Matthaios Olma
MongoDB
matt.olma@mongodb.com

Subhadeep Sarkar
Brandeis University
subhadeep@brandeis.edu

Manos Athanassoulis
Boston University
mathan@bu.edu

ABSTRACT

Search trees, like B$^+$-trees, are often used as index structures in data systems to improve query performance at the cost of index construction and maintenance. For state-of-the-art B$^+$-tree designs used in commercial data systems, this cost is <mark>negligible</mark> if the data arrives as fully sorted on the index attribute. Further, production systems employ a *fast-path* ingestion technique for B$^+$-trees that directly appends the incoming entries to the tail leaf if the data is fully sorted, drastically reducing the index construction cost. However, this is only effective if the incoming data arrives fully sorted or with an extremely small number of out-of-order entries. In addition, the state-of-the-art sortedness-aware design (SWARE) navigates a tradeoff between reads and writes by buffering incoming data to absorb near-sortedness, which comes at the cost of slower query performance and increased overall design complexity.

To address these challenges, we present Quick Insertion Tree (QuIT), a sortedness-aware indexing data structure that improves ingestion performance with minimal design complexity and no read overhead. QuIT maintains in memory a pointer to the *predicted-ordered-leaf* (*poℓe*) that provides a sortedness-aware *fast-path* optimization, and facilitates faster index ingestion. The key benefit comes from accurately predicting *poℓe* throughout data ingestion. Further, QuIT achieves high memory utilization by maintaining tightly packed leaf nodes when the ingested data arrives as near-sorted. This, in turn, helps improve performance during range lookups. Overall, we demonstrate that QuIT outperforms B$^+$-tree (SWARE) by up to 3× (2×) for ingestion, while maintaining the same point lookup performance (up to 1.23× faster). QuIT also accesses up to 2× fewer leaf nodes than the B$^+$-tree during range lookups.

1 INTRODUCTION

Database indexes accelerate query processing by offering fast access to selection predicates. B$^+$-tree indexes [12, 19] are used as the primary index data structure by several popular data systems ranging from relational row-stores [31, 32, 34, 35] like Oracle, SQL Server, PostgreSQL and MySQL to NoSQL systems [13, 20, 30] like MongoDB due to their ability to allow efficient point and range queries. The improved query performance, however, comes at the

Figure 1: (a) QuIT significantly outperforms the existing baselines for any degree of sortedness; and (b) QuIT offers high sortedness-awareness with no additional read penalty while also striking a better right balance between design complexity, tuning, and memory utilization than other baselines.

cost of constructing and maintaining the index as new data is inserted, updated, or deleted from the database [3]. With modern applications requiring data systems to also support faster data ingestion in addition to efficient query processing, the indexing cost becomes prohibitive for workloads with high ingestion rates.

In a B$^+$-tree [12], ingesting a key involves accessing the root note, traversing the tree to find the appropriate leaf node whose range can fit the key, and inserting the key into that node in sorted order. This essentially adds *structure* to the data by establishing a sorted order to the data at the leaf level of the index. The tree traversal, coupled with the required splitting of nodes as the tree grows, heavily consumes the bulk of the indexing cost. There exist techniques that accelerate index ingestion when we have access to the entire dataset *a priori* (e.g., bulk-loading [1, 15]). However, for insertions taking place throughout workload execution, bulk loading is not feasible, and hence, data systems revert to standard ingestion using expensive top-to-bottom tree traversals.

Leveraging Data Sortedness. Our thesis is that since indexing *adds structure* to an otherwise unstructured data collection by creating a fully sorted version of the data, *the indexing effort should be minimal when the data arrives with some intrinsic order* [38]. The key idea here is identifying the correct location (i.e., leaf node), to insert a new key without performing expensive tree traversals. A simple case is when data arrives fully sorted – insertions are always right-deep and are applied only to the tail (right-most) leaf node of the index. Here, maintaining a pointer to the tail leaf suffices to alleviate the indexing effort. In fact, due to its simplicity, this technique is employed in popular commercial systems (e.g., *fast-path* optimization in PostgreSQL [34]). However, the *tail-leaf* optimization is only effective when data arrives fully sorted and fails to capture varying degrees of near-sortedness [37]. A small

number of outliers, equal to the capacity of a single leaf node, is enough to render the tail-leaf optimization useless.

In practice, several applications generate data that is *close-to* being fully sorted (i.e., *near-sorted*). Examples include stock market data, aggregated time-series, intermediate results of joins [5], sorted files that received a small number of updates [5], and correlated columns [2], among others. Recent work on sortedness awareness for indexing (SWARE) [38] aims to exploit partial order in the incoming data by intelligently buffering entries and opportunistically bulk loading them on the fly (into an underlying index, e.g., B$^+$-tree) to accelerate index ingestion. The advantages gained by applying the SWARE paradigm to a B$^+$-tree are illustrated in Fig. 1a, which shows the average ingestion cost when ingesting 500M entries (integer key-value pairs) with varying data sortedness. We benchmark the B$^+$-tree with tail-leaf optimization enabled (henceforth referred to as *tail* or *tail-B$^+$-tree*) the SWARE design (equipped with a 40MB buffer), and our approach. Note that all experiments use the same underlying B$^+$-tree implementation to ensure a fair comparison.

While SWARE outperforms the tail-B$^+$-tree with near-sorted data, it falls back to a B$^+$-tree otherwise. The improved ingestion performance is beneficial for write-heavy workloads, however, accessing the buffer at query time makes SWARE slower for point queries. To reduce the read penalty, SWARE uses a more complex design for the buffer by employing additional data structures (i.e., Zonemaps [29] and Bloom filters [9]), still having up to 26% slower point queries than B$^+$-tree. This requires additional tuning, reduces maintainability, and as a result, hinders adoption.

An ideal tree index should offer high sortedness-awareness through minimal design complexity without hurting query performance.

Simplifying Sortedness-Awareness. In this work, we propose to exploit any intrinsic sortedness in the ingested data to optimize the indexing effort, with the specific goal of incurring minimal additional complexity in the index design. To that end, we propose two fast-path insertion optimizations named *last-insertion-leaf* ($\ell i\ell$) and *predicted-ordered-leaf* ($po\ell e$). The key intuition in both techniques is that they offer a predictor of in-order insertions and direct them to the appropriate leaf to avoid expensive tree traversals.

First, we replace the naïve predictor of *fast-path* insertions from the rightmost (*tail*) leaf to the *last-insertion-leaf* ($\ell i\ell$). This change allows near-sorted ingestion streams to quickly "come back" to the appropriate leaf if we have a small number of unordered entries. While this helps to increase the number of fast-path insertions, inserting an out-of-order entry results in up to two missed fast-path inserts (one top-insert for the unordered entry and another to switch back to the correct leaf node for subsequent entries).

Ideally, we would like to always perform a fast-path insertion, even following an outlier. This is enabled by our second optimization, which is more sophisticated than $\ell i\ell$ yet has minimal design complexity. Specifically, we maintain a pointer to the *predicted-ordered-leaf* ($po\ell e$), which is initiated to the tail leaf. When we receive out-of-order entries, we do not eagerly update $po\ell e$, rather, when the node is split, we decide which of the two resulting nodes should be identified as $po\ell e$ based on the ingested data.

Eliminating Overheads. In the worst case, both techniques easily fall back to regular index ingestion, which is exactly as efficient as the underlying B$^+$-tree. This way, we avoid a fraction of tree traversals without incurring any other overhead. The required metadata is also minimal: one pointer to the fast-path ($\ell i\ell$ or $po\ell e$), the smallest and largest values that the fast-path node can accept, and only for $po\ell e$, the size and the smallest key of the previous leaf.

Quick Insertion Tree. We propose *Quick Insertion Tree* (QuIT), a lightweight indexing data structure that supports fast data ingestion using the $po\ell e$ fast-path optimization. QuIT adapts to the sortedness of the incoming data to facilitate fast index appends. In the ingestion experiment shown in Figure 1a, QuIT outperforms the tail-B$^+$-tree by up to $\approx 2.5\times$, and the SWARE design by $\approx 2\times$.

Since $po\ell e$ is identified as a node receiving ordered entries, we use an *In-order Key estimatoR* (*IKR*) to guide its update policy (inspired by the inter-quartile range [14]). Using *IKR* enables QuIT to employ a variable split factor that better packs in-order entries in its leaf nodes, in addition to redistributing entries between underutilized leaf nodes. This helps the index to improve both space utilization when ingesting near-sorted data and read performance during range lookups. Further, point lookups in QuIT are similar to B$^+$-tree, hence, *QuIT ingestion benefits come with no read penalty.*

Figure 1b shows a qualitative comparison between tail-B$^+$-tree, SWARE, and QuIT based on sortedness-awareness (ingestion benefits), read cost, complexity, tuning, and memory utilization. Overall, QuIT outperforms both the tail-leaf optimization found in production systems and the SWARE paradigm for sortedness awareness without incurring any additional read penalty. QuIT achieves this with little to no tuning, minimal design complexity, and improved memory utilization. Further, QuIT employs classical concurrency strategies to allow for multi-threaded execution with some additional care for the metadata used. Overall, QuIT's lightweight design allows for easy adoption into production data systems.

Contributions. Our work offers the following contributions:

- We propose two simple, yet powerful *fast-path* optimizations for B$^+$-trees: $\ell i\ell$ that *reuses* the *last-insertion-leaf* to avoid a full tree traversal, and $po\ell e$ that exploits near-sortedness to predict which leaf should receive the next in-order insertion (§3).
- We present an *In-order Key Estimator* (§4.1) that updates $po\ell e$ (§4.2). *IKR* also allows for variable-split ratio and redistribution in leaf nodes to ensure higher space utilization (§4.3).
- We integrate the above techniques with minimal metadata and tuning (§4.4) and support for concurrent execution (§4.5) into the Quick Insertion Tree(or QuIT), a general-purpose index that supports sortedness-aware fast ingestion.
- We extensively evaluate QuIT and its core components (§5). We show that QuIT significantly outperforms a state-of-the-art B$^+$-tree by up to 2.5× and SWARE by up to 2× during near-sorted data ingestion. QuIT is also 1.26× faster than SWARE during point lookups and improves its memory footprint by up to 49% when compared to B$^+$-tree.
- Finally, we demonstrate that QuIT scales well with data size and with concurrent execution.

2 BACKGROUND AND MOTIVATION

In this section, we provide the necessary background to data sortedness and the associated index construction and maintenance cost. We discuss the existing optimization techniques that attempt to exploit sortedness as a resource, and why they fall short.
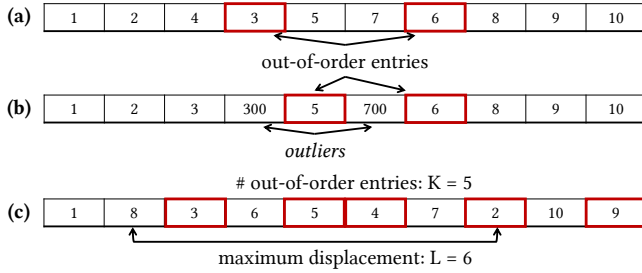
Figure 2: Examples of (a) out-of-order entries and (b) outliers in a data collection. (c) The $K-L$-sortedness metric effectively captures the number of out-of-order entries as K and the maximum displacement of out-of-order entries as L.

Quantifying Data Sortedness. Data *sortedness* captures the difference between the arrival order and indexed order of data over the indexed attribute. Several metrics have been proposed in the literature that aim to quantify the sortedness of a stream of data [5, 11, 24, 28]. One way to quantify data sortedness is to count the number of entries that are out of order in a dataset [28]. By this quantification, a *fully sorted* data collection has no out-of-order entries, whereas a *scrambled* data collection has all or nearly all of its entries out of order. A *nearly-sorted* data collection has only a few out-of-order entries in an otherwise sorted data collection. In general, out-of-order entries are identified as those that are smaller than their preceding key in a monotonically increasing data stream, as shown in Fig. 2a, or vice versa. Further, entries that deviate considerably from the overall expected value in a near-sorted data stream and that may (or may not) be in order with respect to their preceding entry are categorized as *outliers*. For example, in Fig. 2b, although the entries 300 and 700 are in order with their respective preceding keys, they are considered *outliers* as they deviate significantly in magnitude with neighboring entries. Note that outliers are easy to identify when the data set is available in its entirety, however, identifying them accurately in an incoming data stream is challenging.

The K−L-Sortedness Metric. The $K-L$-sortedness metric [37] inspired by Ben-Moshe et. al. [5] more comprehensively quantifies data sortedness by accounting for both *the out-of-order entries* and *the distance by which they are out of order*. The number of out-of-order entries is denoted by K, and the maximum displacement of an unordered entry from its in-order position is denoted by L. An example of a nearly sorted data collection based on $K-L$-sortedness is shown in Fig. 2c where K=5 entries are out of place and the out-of-place entries are displaced by at most of L=6 index positions.

Sortedness-Aware Indexes. Raman *et al.* [38] proposed the SWARE indexing paradigm that captures the sortedness of a data stream using an in-memory buffer and opportunistically bulk load (on-the-fly) incoming data to the underlying tree index (e.g., B+-tree). This helps improve the ingestion performance for near-sorted or even less-sorted data streams, as shown in Figure 1a.

The benefits gained by buffering entries during ingestion come at the expense of query performance, as every query now has to first search the buffer. SWARE partially addresses this overhead by employing auxiliary data structures like Zonemaps [29] and Bloom filters [9], in addition to a query-driven partial-sorting technique


Figure 3: The tail-leaf optimization is effective only for an extremely high degree of sortedness and leads to no fast-inserts when 1% or of the entries are out of order.

that is inspired by Cracking [21, 22]. Yet, point queries in SWARE are up to 26% slower than the baseline B+-tree, a cost that becomes prohibitive as the fraction of reads in the workload increases.

More importantly, adding Zonemaps and Bloom filters to the design implies that insertions to the index are no longer simple appends to the buffer. Every insert first checks if it is arriving in-order to the preceding key, and if otherwise, performs a linear scan of the Zonemaps to identify overlapping pages within the buffer. Additionally, the inserted key is also indexed through a couple of layers of Bloom filters that require re-calibration during every buffer flush. The memory buffer and the metadata (including the auxiliary data structures) also increase the memory footprint (e.g., for indexing 1 TB of data, the memory requirement can be more than 10 GB). Thus, in addition to imposing a penalty on the lookup performance, SWARE also requires careful tuning to ensure that the benefits of buffering outweigh tree traversals during regular insertions to guarantee overall performance improvement.

Tail-Leaf Insertion. Tail-leaf insertion in B+-trees is a simple and effective *fast-path* optimization that benefits from incremental in-order ingestion to the index. The tail-leaf *fast-path* essentially maintains one additional pointer to the rightmost leaf (tail) of the index along with its smallest allowed value. Any newly ingested key that is greater or equal to that value is directly inserted into the tail-leaf that is naturally cached, rather than traversing the tree. While this optimization is rarely useful when data does not arrive in sorted order, surprisingly, this fast-path optimization fails to accelerate index ingestion even when data arrives near-sorted. As soon as the number of outliers inserted exceeds one node worth of data, the tail-leaf only contains outliers, resulting in a "stale" fast-path. This results in future insertions (even for near-sorted data) reverting to traditional incremental ingestion to the index (referred to as *top-inserts*) and missing the opportunity to utilize the fast-path (referred to as *fast-inserts*).

Tail-leaf is Only Helpful for Extremely High Sortedness. Fig. 3 shows the fraction of fast-inserts when ingesting 5M integers into a tail-B+-tree, as we vary data sortedness. Specifically, we vary the fraction of out-of-order entries, which are positioned uniformly and randomly in the workload. As expected, the tail-leaf optimization is effective for sorted data (i.e., 0% out-of-order entries) or extremely near-sorted data (i.e., very few out-of-order entries). Thus, while we get negligible top-inserts for 0.01% out-of-order entries, the tail-leaf optimization's efficiency drops to only 23% (11%) fast-inserts for 0.05% (0.1%) out-of-order entries, and, ultimately, to less than 1% fast-inserts for 1% out-of-order entries and beyond. This renders the
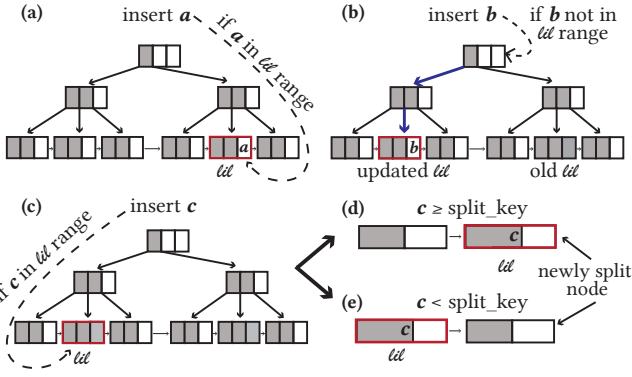
Figure 4: Fast-path ingestion using $\ell i\ell$: (a) a newly inserted key is added to $\ell i\ell$ if within $\ell i\ell$-range; (b) a top-insert updates $\ell i\ell$ pointer to the leaf node where we insert the key; (c) when an insert to $\ell i\ell$ causes a split, (d) $\ell i\ell$ is updated if $c \geq split\_key$; otherwise (e) $\ell i\ell$ stays as is.

tail-leaf optimization impractical for most near-sorted workloads as *fast-path* ingestion is very rarely used.

3 REINVENTING FAST-PATH OPTIMIZATION

We now propose *last-insertion-leaf* ($\ell i\ell$), a renewed fast-path optimization technique for B$^+$-trees that offers superior ingestion performance for near-sorted workloads, and lay the groundwork for *predicted-ordered-leaf* ($po\ell e$), QuIT's key ingredient.

Tracing the Leaf of the Last Insertion. Contrary to the tail-leaf optimization, we maintain a pointer to the *last-insertion-leaf* or $\ell i\ell$ for short. At any point during a workload execution, $\ell i\ell$ points to the leaf node to which the most recent entry was inserted. A subsequent insert may be added to the $\ell i\ell$-node if it falls within its range as shown in Fig. 4a. Otherwise, we revert to a *top-insert* followed by an update to the $\ell i\ell$-pointer (Fig. 4b). The $\ell i\ell$-pointer is also updated if a newly inserted entry results in splitting the $\ell i\ell$-node (Fig. 4c). In this case, we update the $\ell i\ell$-pointer if the inserted key is placed into the newly created node from the split (Fig. 4d), or keep it unchanged otherwise (Fig. 4e).

Modeling $\ell i\ell$ Benefits. We quantify the expected efficiency of $\ell i\ell$ by estimating the fraction of data that would be fast-inserted into the tree as a function of the out-of-order entries. Ideally, a sortedness-aware index would fast-insert all in-order entries and perform top-inserts only for entries that are out of order. However, $\ell i\ell$ performs a top-insert both when (a) an out-of-order entry follows an in-order entry (outlier or not) and (b) an in-order entry follows an outlier. Conversely, $\ell i\ell$ would only succeed when we have two in-order entries in a row. We calculate the probability of two consecutive entries being in order as follows. We assume that we insert n entries into a B$^+$-tree and that a fraction, k, of those entries are out of order. Then, the number of in-order entries, y, is: $y = n \cdot (1 - k)$, and the probability of a fast-insert (which happens when two consecutive entries are in-order), FI, is given as:

$$FI = \frac{y}{n} \cdot \frac{y-1}{n-1} \underset{\text{for large } n}{\approx} \left(\frac{y}{n}\right)^2 = (1-k)^2 \qquad (1)$$


Figure 5: (a) The *last-insertion-leaf* optimization significantly outperforms tail-leaf insertions in B$^+$-trees for highly sorted data. (b) Simulation of the expected fraction of top-inserts using the $\ell i\ell$-pointer while varying data sortedness.

Evaluating $\ell i\ell$. While the tail-leaf optimization results in virtually no fast-inserts even with only 1% out-of-order entries, Eq. (1) shows that $\ell i\ell$ should manage to achieve 98% fast-inserts in the same workload, which is corroborated experimentally. Fig. 5a shows the fraction of fast-inserts (on the y-axis) when ingesting $5M$ entries (integer K-V pairs) as we vary the fraction of out-of-order entries k in the x-axis. Firstly, for fully sorted data both the tail and $\ell i\ell$ optimizations invoke their respective fast-path insertion routines and avoid expensive top-inserts altogether. We observe that $\ell i\ell$ is indeed able to perform 98% fast-inserts for a workload with $k = 1\%$, 90% fast-inserts for a workload with $k = 5\%$ while also performing around almost no fast-inserts for $k = 100\%$. The very few fast-inserts for the latter are due to two consecutive out-of-order entries targeting with a very low probability the same node, which is slightly higher initially when the tree is small.

The superior benefits of $\ell i\ell$ against the tail-leaf optimization is because the latter works well only when an incoming entry can be correctly positioned in the tail leaf, the probability of which is very low. On the other hand, we observe a gradual decrease of fast-inserts performed in $\ell i\ell$ as we decrease data sortedness. While $\ell i\ell$ initially points to the tail leaf, it updates to the last insertion leaf as soon as we encounter an entry that is top-inserted elsewhere. This allows all subsequent in-order entries to be ingested through the fast path to the new $\ell i\ell$-node or revert to the "correct" leaf node if $\ell i\ell$ points to a leaf filled with outliers.

Headroom of Improving $\ell i\ell$. While $\ell i\ell$'s design offers the opportunity to perform better than the state-of-the-art tail-leaf optimization, there is still a lot of room for improvement. To quantify the headroom of improvement, we compare the expected number of fast-inserts performed by $\ell i\ell$ with the theoretical ideal design of a sortedness-aware fast-path optimization. We use Eq. (1) to estimate the number of fast-inserts as a fraction of the workload, while the ideal is all in-order entries. Fig. 5b, shows that while $\ell i\ell$ (solid black line) is expected to clearly outperform tail-leaf insertion (in dashed black line) – from the experiment in Fig. 5a, the number of fast-inserts quickly drops as the number of out-of-order entries increases. Instead, for an ideal sortedness-aware index (green line with a triangle marker), the fraction of fast inserts are expected to be linearly proportional to the in-order entries. The area between the solid black and green lines presents the headroom for improvement.

Focusing on the missed opportunity, $\ell i\ell$ performs two top-inserts per out-of-order entry: (i) one for an out-of-order entry, and (ii) one for the in-order after $\ell i\ell$ is updated to the wrong leaf.

Optimal sortedness-awareness should incur, at most, one top-insert per out-of-order entry.

Thus, we expect that the number of top-inserts (that can be 3-4× more expensive than fast inserts depending on the height of the tree) will be halved by such an ideal design, reducing substantially the overall cost of near-sorted data ingestion. Finally, when inserting near-sorted data we can improve the space utilization of the index by better packing leaf nodes with in-order entries and using a variable split ratio [38]. On the other hand, for tail-leaf optimization and $\ell i\ell$, the higher the data sortedness the lower the space utilization, since every node split will leave a half-full node that will never receive any future insert.

Next, we propose a new design that bridges the aforementioned performance headroom with a more robust fast-path optimization, while also offering better space utilization.

4 QUICK INSERTION TREE

We present *Quick Insertion Tree* (QuIT), an indexing data structure that is sortedness-aware *by design* and offers superior ingestion performance along with better space utilization when ingesting near-sorted data. At its foundation, QuIT is similar to B$^+$-tree – its root and internal nodes contain a list of keys and pointers, while the leaf nodes contain the data entries. However, QuIT employs a sortedness-aware fast-path optimization that benefits ingestion workloads that have at least some degree of intrinsic data sortedness. If the data is completely scrambled, QuIT effectively behaves like B$^+$-tree for writes. The key advantage of QuIT over other sortedness-aware counterparts [38] is the lack of any additional read penalty when compared to a B$^+$-tree. It also has minimal design complexity and requires little to no tuning. Through the rest of this section, we present the architecture of the Quick Insertion Tree.

4.1 A Robust Fast-Path Optimization

Predicting the Ordered Leaf. In Section 3 (Fig. 5b), we pointed out that an ideal sortedness-aware index would perform expensive top-inserts only for entries that are out of order, while all other entries should be ingested using the fast path. The fundamental limitation of $\ell i\ell$ is that it naïvely switches the fast-path access pointer (i.e., the $\ell i\ell$-pointer) based on the most recent insert, even if the entry ingested is out of order. We address this by replacing $\ell i\ell$ with a new leaf node pointer to the *predicted-ordered-leaf* ($po\ell e$) node, that tracks the leaf that is most likely to accept the future in-order entries. Similarly to $\ell i\ell$, out-of-order entries are top-inserted. However, unlike $\ell i\ell$, the pointer to $po\ell e$ may be updated only when $po\ell e$ splits. The newly created node from the split will be identified as $po\ell e$ if its smallest key is not an *outlier*, while $po\ell e$ remains unchanged otherwise. To guide this update policy, we build a lightweight outlier predictor called *In-order Key estimatoR* (IKR).

Identifying Outliers. Assume we are inserting entries into the index with keys following an increasing order. Let $po\ell e_{size}$ and $po\ell e\_prev_{size}$ denote the respective number of entries in $po\ell e$ and its preceding node $po\ell e\_prev$. We want to identify if $po\ell e$ needs to be updated upon split. Now, let p and q be the values of the smallest key in the $po\ell e\_prev$ and $po\ell e$, as shown in Fig. 6a. Since $po\ell e$ contains in-order entries, q is not an outlier, and necessarily p is also not an outlier as it precedes q. When splitting, we would

**Figure 6: Updating $po\ell e$: (a) post splitting $po\ell e$, we use the smallest key (r) in the newly split node ($po\ell e\_next$) and compare to an estimation (x) from the *IKR*; (b) if $r \leq x$, we update $po\ell e$; otherwise, (c) if $r > x$ we leave $po\ell e$ as is.**

Algorithm 1: Updating *predicted-ordered-leaf*

Data: $p = po\ell e\_prev_{min}$, $q = po\ell e_{min}$, $entry = (key, value)$
Init: $scale = 1.5$

1 **if** $q \leq key < po\ell e_{max}$ **then**   // fast-insert
2   **if** $po\ell e$ *is full* **then**
3      $po\ell e\_next \leftarrow po\ell e.split()$;   // return new leaf
4      $e \leftarrow po\ell e\_next_{min}$;
5      $x \leftarrow q + \left(\frac{q-p}{po\ell e\_prev_{size}}\right) \cdot po\ell e_{size} \cdot scale$;
6   **if** $e \leq x$ **then**
7        $po\ell e\_prev \leftarrow po\ell e$;
8        $po\ell e \leftarrow po\ell e\_next$;
9    $po\ell e.insert(entry)$;
10 **else**
11    $l_t \leftarrow top\_insert(entry)$; // return leaf that accepts
12   **if** $l_t = po\ell e\_next$ **then**   // *$po\ell e$* catches up
13      $po\ell e\_prev \leftarrow po\ell e$;
14      $po\ell e \leftarrow po\ell e\_next$;

---

like to identify whether the smallest key in the new node created from the split, r (i.e., the split key), is an outlier. This helps decide whether to keep the pointer to $po\ell e$ unchanged (if r is an outlier) or move $po\ell e$ to the new node (if r is not an outlier).

Our lightweight *IKR* estimator (inspired by *Interquartile Range* outlier detection [14]) calculates the maximum acceptable domain for a non-outlier key. Any key beyond this range is considered an outlier (denoted by x) as follows:

$$x = q + \left(\frac{q-p}{po\ell e\_prev_{size}}\right) \cdot po\ell e_{size} \cdot scale \qquad (2)$$

The term $\frac{q-p}{po\ell e\_prev_{size}}$ in Eq. (2) calculates the density between p and q (two non-outliers). To ensure enough data for prediction, we bound $po\ell e\_prev_{size} \geq 50\%$ (which is always true in traditional B$^+$-tree-node-splitting). The *scale* allows a small buffer to capture small deviations in density that are inherent in the data. Following standard practice [14], we use $scale = 1.5$. Consequently, from Eq. (2), any $key > x$ is considered an outlier.

4.2 Fast-Path Insertion in $po\ell e$

We now describe the application of $po\ell e$ as a fast-path ingestion technique in the Quick Insertion Tree and outline the steady-state insertion algorithm in Algorithm 1.

Algorithm 2: Variable split strategy

Data: $q = pole_{min}$, $entry = (key, value)$

Init: $scale = 1.5$, $def\_split\_pos = \frac{leaf_{capacity}}{2}$

1 if $leaf \neq pole$ **then**
2 $leaf.split(def\_split\_pos);$ `// split leaf at 50%`
3 else if $pole\_prev_{size} \geq def\_split\_pos$ **then**   `// use IKR`
4 $l \leftarrow leaf.position(q + \left(\frac{q-p}{pole\_prev_{size}}\right) \cdot scale);$
5   **if** $l > def\_split\_pos$ **then**
       `// take one non-outlier to new node`
6 $pole\_next \leftarrow pole.split(l - 1);$
7 $pole\_prev \leftarrow pole;$
       $pole \leftarrow pole\_next$
8   **else**
       `// move all outliers to new node`
       $pole\_next \leftarrow lol.split(l);$
9 else
    `// redistribute entries from pole_prev`
10 $lol.redistribute(pole\_prev);$

Figure 7: The default split position is in the middle of $pole$ **(50%). When** $pole\_prev_{size} \geq 50\%$**, we use** *IKR* **to identify outliers. (a) If outliers occupy** $< 50\%$ **of** $pole$ **(**$l > def\_split\_pos$**), we split at** $l - 1$**, moving a non-outlier value to the newly split node and update** $pole$**; (b) Otherwise, we split at** l **(using the default** *IKR***) and keep** $pole$**-pointer as is; (c) When** $pole\_prev_{size} < 50\%$**, we redistribute entries to** $pole\_prev$**, until** $pole\_prev_{size} = 50\%$**.**

Initialization. The initial state of the index is represented by a single leaf node in the tree that is also its root. We also mark this leaf as $pole$. When this leaf first splits, we create a new root node and add the pivot pointers to its two child nodes (similar to a B$^+$-tree). We mark the leaf that received the latest insert as the $pole$-node.

Steady State. In steady-state, an entry inserted to the index utilizes $pole$ if its key is within the range of $pole_{min}$ and $pole_{max}$, i.e., the smallest and largest keys that can be insrted into $pole$. When $pole$ is also the tail-leaf, we omit the upper bound check on $pole_{max}$.

Splitting the $pole$**-node.** Once $pole$ is full, we split it and update the pointer to $pole$ as shown in Fig. 6. We refer to the newly created node from the split as $pole\_next$. We compare x from Eq. (2) with the smallest key in $pole\_next$, denoted by r. We update $pole$ to $pole\_next$ if $r \leq x$ (Fig. 6b), and leave it as-is otherwise (Fig. 6c).

Catching Up to Predicted Outliers. When $r > x$ (i.e., when r is an outlier), we infer that all entries of the new node are outliers. So, any future in-order entries will belong to the node that was split, and thus, we do not update $pole$ after splitting. Eventually, if keys are largely in-order, entries in the $pole$-node may catch up to previously marked outliers in $pole\_next$. Hence, when an entry is top-inserted to $pole\_next$, we check if it is still an outlier using *IKR* and, only if not, update $pole$ to $pole\_next$.

4.3 Improving Space Utilization

The $pole$ optimization offers a robust and sortedness-aware fast-path access to the leaf level of the index that avoids tree traversals for in-order data to boost ingestion performance. However, the current design retains the same worst-case space utilization as a B$^+$-tree when ingesting fully-sorted data. That is, due to consecutive right-deep insertions, every node will be exactly half full, wasting 50% of space. Ideally, a sortedness-aware index should exploit the underlying data sortedness to improve its space utilization.

Finding Better Split-Points. QuIT exploits *IKR* and $pole$ to further improve space utilization. First, in Algorithm 2, we detail a variable

split strategy in the leaf nodes through which ordered data can be more tightly packed. We re-use *IKR* to identify the outliers and determine the optimal split-points for the node. Note that this is fundamentally different from Algorithm 1, where we first split $pole$ by default at 50% and only use the *IKR* to update the $pole$-pointer. Below, we discuss the major decisions in Algorithm 2.

Splitting a Node that is Not the $pole$**-node.** Similar to a classical B$^+$-tree, we split a leaf node that is not $pole$ at 50% ($def\_split\_pos = leaf_{capacity}/2$ in lines 1-2).

Splitting $pole$ **When** $pole\_prev$ **is At Least Half-full.** When $pole$ is about to split and its preceding node ($pole\_prev$) is at least half-full, we use *IKR* to safely identify the split position of an in-order key. We identify the position (l) of the first key greater than the estimated acceptable value lower bound in $pole$ (line 4).

When $pole$ **Contains Only Few Outliers.** If $l > def\_split\_pos$ (i.e., if $pole$ mostly contains in-order entries), we split the node (and update it as $pole$) at $l - 1$ (lines 7-9), taking one in-order value to the newly created node from the split, as shown in Fig. 7a. This ensures that the split leaf is at least half full while the new $pole$ has more space to accommodate future fast-insertions.

When $pole$ **Contains Mostly Outliers.** When $l \leq def\_split\_pos$, $pole$ contains mostly outliers. In this case, we split $pole$ at l, and move all outliers to the newly created node from the split (as shown in Fig. 7b). We leave the pointer to $pole$ unchanged as it now has enough space for future fast-insertions.

Redistribution When $pole\_prev$ **is Less than Half-full.** In case that at $pole$-splitting time, $pole\_prev$ is less than half full (e.g., due to an earlier variable split), using *IKR* may lead to an inaccurate estimation as it does not have enough data. Instead, we redistribute entries (line 10) from $pole$ to $pole\_prev$ until the latter is exactly half full (shown in Fig. 7c). This allows future splits in the $pole$ node to use the variable split strategy.

Resetting Fast-Path When $pole$ **Goes Stale.** As with any other fast-path optimization, there exists a scenario where $pole$ becomes

Index	B$^+$-tree	tail-B$^+$-tree	$\ell i\ell$-B$^+$-tree	QuIT
root_id	✓	✓	✓	✓
head_id	✓	✓	✓	✓
tail_id	✓	✓	✓	✓
fp_path[]		✓	✓	✓
fp_size		✓	✓	✓
fp_min		✓	✓	✓
fp_max			✓	✓
fp_id			✓	✓
$po\ell e$_prev_size				✓
$po\ell e$_prev_min				✓
$po\ell e$_prev_id				✓
$po\ell e$_fails				✓

Table 1: Metadata used by different indexes. QuIT needs less than 20 bytes of additional metadata

stale due to certain workload characteristics, leading to unexpectedly many top-inserts. For example, this can be caused by workloads that do not exhibit any sortedness, or by corner cases of specific sequences of in-order and out-of-order insertions that may throw *IKR* off. We recover from the *stale* state by *resetting $po\ell e$* to the leaf that accepted the latest insert. However, unlike $\ell i\ell$, we do not adjust $po\ell e$ for every top-insert. Rather, we do so only if we have already performed a number of consecutive top-inserts. We set this threshold as $T_R = \lfloor \sqrt{leaf_{capacity}} \rfloor$ to offer a balanced outcome.

Metadata Digest. We summarize the overall metadata used by the B$^+$-tree and its variants that use a fast-path optimization, namely, tail and $\ell i\ell$, as well as QuIT, in Table 1. We denote by the prefix fp any metadata common to all fast-path optimization. Note that fp_path (inspired by PostgreSQL [34]) stores the root-to-leaf path to the fast-path node. This allows efficient access to internal nodes that are ancestors of the fast-path node during splits. We highlight that QuIT requires only a few bytes of additional state *for the entire tree* to offer sortedness-aware ingestion with no read penalty.

4.4 Other QuIT Operations

Lookups. Lookups in QuIT are identical to the classical B$^+$-tree. A point lookup-path starts at the root node and uses key comparisons to follow the pivot pointers, to navigate to the appropriate leaf node that may contain the target entry. A binary search on the keys in the leaf node returns whether the search key is present in the index or not. A range lookup on keys in the range $[start, end]$, first performs a point lookup on $start$ to locate the first key $\geq start$. It then uses the interlinked pointers in the leaf node to scan entries across leaf nodes until an entry $\geq end$ is found.

Deletes. Delete operations in QuIT are also exactly the same as the B$^+$-trees. Deletion of an entry begins with a point lookup on the target key. If the key is found, in QuIT, we typically remove the key from the corresponding leaf node and apply rebalancing to ensure that the leaf node and all internal nodes leading to the leaf node are at least half full. Only deletes of entries in the $po\ell e$-node do not rebalance eagerly. In case the key marked to delete is the only key in $po\ell e$, we reset $po\ell e$ to $po\ell e$_prev.

4.5 Concurrency Control

Concurrency control protocols that have been widely studied for B$^+$-trees [4, 7] can be applied out-of-the-box to QuIT.

Locking Protocol for Ingestion. In B$^+$-trees, the internal nodes simply redirect searches to the next level that contains either an internal node or a leaf node. Hence, completely and exclusively locking the entire path is wasteful. We employ a simple protocol inspired by classical lock-crabbing [18].

Classical lock-crabbing [18] starts at the root node and descends down the tree, acquiring exclusive locks at every node along the insertion path. If a node along the path is not full, an insertion does not result in a split. Therefore, all the acquired locks in the preceding levels of the path are released. A lock on the entire path is acquired only when every node along the path is full. The locks are retained until the insertion completes, as a split in the lowest level of the tree (i.e., the leaf node) can potentially propagate to every internal node along the path, including the root.

Every insert in QuIT first acquires a lock on the fast-path metadata to verify if a $po\ell e$ insertion can be utilized. If the key is within the range of the fast-path access, and $po\ell e$ is not full, an exclusive lock of the $po\ell e$ leaf is acquired and released post-insertion. When $po\ell e$ is full, we acquire locks to the preceding levels of the index (i.e., the internal nodes), by utilizing the lock-crabbing protocol on fp_path. Additionally, a split would either result in the creation of a new node or a redistribution of entries from $po\ell e$ to $po\ell e$_prev. In the former case, we only acquire a lock on the metadata of $po\ell e$_prev (i.e., $po\ell e$_prev_id, $po\ell e$_prev_min, and $po\ell e$_prev_size). For redistribution, we also acquire a lock on $po\ell e$_prev and its associated metadata. For a top-insert, we utilize lock-crabbing, as in a B$^+$-tree.

Locking Protocol for Lookups. Concurrent lookups in QuIT essentially follow the same procedure as in a B$^+$-tree with lock-crabbing. We start from the root node, albeit, acquiring shared locks on every internal node along the access path. Once a shared lock on a node is obtained, we immediately unlock its parent. Range lookups similarly obtain a lock on the first leaf node that is accessed by the query. However, we also acquire shared locks on subsequent leaf nodes that are accessed through the interlinked pointers.

5 EVALUATION

We now show the benefits of QuIT by comparing it to (i) a textbook B$^+$-tree (we call this a *classical B$^+$-tree*) that only performs top-inserts, (ii) a B$^+$-tree with tail-leaf insertion enabled (we call this a *tail-B$^+$-tree*), (iii) a B$^+$-tree with the *last-insertion-leaf* optimization (we call this a $\ell i\ell$-B$^+$-tree), and the state-of-the-art sortedness-aware index design, SWARE [38]. We test using various configurations, varying data sortedness and data size, stress-testing the fast-path, and we also experiment with real-world data.

Experimental Setup. We run experiments using our in-house server with 128GB of DDR5 main memory at 4800 MHz and a 1.9TB NVMe SSD. The server runs Rocky Linux (version 9.3) and is equipped with two sockets of Intel Xeon Gold (6442Y) 2.6 GHz processors (24 cores), each capable of supporting 48 threads.

Index Design and Default Setup. We use an in-memory implementation of the B$^+$-tree (inspired by state-of-the-art [8]), and implement tail, $\ell i\ell$, and $po\ell e$ optimizations on this platform. We also

Aneesh Raman, Konstantinos Karatsenidis, Shaolin Xie, Matthaios Olma, Subhadeep Sarkar, and Manos Athanassoulis

Figure 8: QuIT outperforms all current baselines for any degree of data sortedness during ingestion.

Figure 9: QuIT considerably reduces the top-insertions by using the $pole$ optimization for any degree of sortedness.

implement QuIT on the same platform that includes features such as the variable-split, redistribute, and reset strategies. We use the same B$^+$-tree implementation as the underlying index when comparing with the SWARE paradigm by extending the API to support bulk loading. For the SWARE buffer, we deploy the open-sourced code[1] out-of-the-box and default to a buffer size equivalent to 1% of the total data size. The default entry size in all our experiments is 8B (with 4B keys), and we use a 4KB page size that fits up to 510 entries in the leaf nodes. We set $T_R = \lfloor \sqrt{leaf_{capacity}} \rfloor = \lfloor \sqrt{510} \rfloor = 22$ to trigger a reset of the $pole$ fast-path in QuIT. We make our code available on GitHub at https://github.com/BU-DiSC/quick-insertion-tree.

Workloads. To evaluate QuIT with varying degrees of sortedness, using the workload generator from the Benchmark on Data Sortedness (BoDS) [37]. BoDS uses the $K - L$-sortedness metric to create a family of differently sorted data collections. The generator takes as arguments (i) the number of entries (N) in the data collection, (ii) the number of unordered entries (K) and (iii) the maximum displacement (L) both as a fraction of the total data size, (iv) the distribution of data sortedness (α, β), and (v) a seed value.

Default Workload. For our experiments, we set N=500M to generate datasets of size 4GB. The skew-parameters α and β are set to 1 to uniformly distribute the unordered entries. Our query workloads (executed post data ingestion) contain 5M (1% of total data size) point lookups that are generated uniformly and randomly on existing keys in the index. Further, our query workloads also contain 1000 range lookups that are generated for random ranges in the key domain with three levels of selectivity: 0.1%, 1%, and 10%.

5.1 Benefits of Quick Insertion Tree

We first demonstrate the benefits of QuIT against three baselines – classical B$^+$-tree, tail-B$^+$-tree, and $\ell i \ell$-B$^+$-tree for ingestion, memory utilization and query performance. In these experiments, we vary the fraction of out-of-order entries (K) in the data collection while setting their maximum displacement (L) to 100%.

QuIT Outperforms the State of the Art in Insert Performance. Fig. 8 shows the speedup (y-axis) during ingestion for tail-B$^+$-tree, $\ell i \ell$-B$^+$-tree, and QuIT relative to the classical B$^+$-tree when varying data sortedness (x-axis). QuIT significantly outperforms the tail-B$^+$-tree and the classical B$^+$-tree in terms of raw data ingestion performance. For fully sorted data, both QuIT and a tail-B$^+$-tree

offer $\approx$ 3$\times$ better performance than a classical B$^+$-tree, as entries are directly inserted to the tail leaf using the fast-path optimization, rather than performing tree traversals. However, the performance of the tail-B$^+$-tree degrades very quickly as the data becomes even slightly unsorted, making it comparable to that of a classical B$^+$-tree. QuIT, on the other hand, exploits any inherent data sortedness and offers up to $\approx$ 2.5$\times$ better performance for near-sorted data ($K < 25\%$). Even for less-sorted data (K=25%), QuIT is still 1.4$\times$ better than both the classical B$^+$-tree and tail-B$^+$-tree, as both the baselines always perform top-inserts, accumulating a redundant indexing cost. The tail-B$^+$-tree performs better than the classical B$^+$-tree only with fully sorted data by allowing fast-path ingestion to the right-most leaf node in the index. This fast-path, however, becomes stale with even near-sorted data ($K \geq 1\%$), leading to performance degradation as its window to capture unorderedness in the data is too narrow (limited to the tail-leaf's range). Meanwhile, QuIT carefully adapts its fast-path (i.e., $pole$) to the data sortedness, thereby, delivering better performance.

QuIT Offers a More Robust Performance than $\ell i \ell$-B$^+$-tree. In Fig. 8, we also observe that QuIT performs up to 10% better during ingestion than the $\ell i \ell$-B$^+$-tree. The $\ell i \ell$ optimization is simple and effective in improving ingestion performance over a tail-B$^+$-tree, and in fact, offers comparable performance to QuIT when the data is very nearly sorted (K<5%). However, as the data becomes less sorted, QuIT's benefits become more pronounced and is significantly able to outperform the $\ell i \ell$-B$^+$-tree (and all other baselines) due to its robust fast-path optimization.

This can be explained by the trend we observe in Fig. 9 that shows the fraction of insertions that utilize the fast-path (y-axis) as we vary data sortedness (x-axis). Here, we omit the classical B$^+$-tree, since it only performs top-inserts. The tail-B$^+$-tree (grey bar with vertical lines) performs fast-inserts only for fully sorted data, and even a minor increase in the unorderedness causes a steep degradation in its ability to utilize the fast-path, thereby, reducing it to a classical B$^+$-tree. The $\ell i \ell$-B$^+$-tree performs significantly more fast-inserts, limiting top-inserts to only $\approx$ 35% of the total insertions even when the data is almost scrambled (K=50%). This matches the expectation from Eq. (1) presented earlier in §3. Meanwhile, QuIT is the most efficient, performing approximately only as many top-inserts as there are out-of-order entries, and very closely resembles the optimal sortedness-aware index behavior discussed in Fig. 5b. This is because $\ell i \ell$ simply updates its fast-path access

[1]SWARE code-base: https://github.com/BU-DiSC/sware

Figure 10: (a) Variable split factor in QuIT allows for higher occupancy in the leaf nodes; (b) Point lookups do not incur any read overhead; (c) Range queries are faster as they access fewer nodes during lookups.

Index	% unordered entries							
	0%	1%	3%	5%	10%	25%	50%	100%
QuIT	1.96×	1.5×	1.41×	1.32×	1.16×	1.09×	1.01×	1×

Table 2: Space reduction of QuIT over all B$^+$-tree baselines. $\ell i\ell$-B$^+$-tree and tail-B$^+$-tree are omitted because they have the same memory footprint as the basic B$^+$-tree.

pointer based on the last insertion to the index (even during top-inserts). Meanwhile, QuIT has a more robust maintenance scheme that utilizes $po\ell e$ as well as its reset strategies that help effectively and efficiently predict the fast-path to inserting the data.

QuIT Improves the Overall Memory Utilization. The *IKR*-based variable node split strategy in QuIT is crucial to reduce the overall memory footprint of the index. Table 2 compares the normalized memory footprint of the indexes, the lower signifying the better. The memory utilization of the tail-B$^+$-tree and the $\ell i\ell$-B$^+$-tree is identical to that of classical B$^+$-tree, as they all split a full node in half. We observe that QuIT's memory footprint is up to 1.96× smaller due to better space utilization on the leaf nodes, as shown in Fig. 10a. Here, we vary the data sortedness on the x-axis and show the average leaf node occupancy (i.e., # entries in a leaf node) as a fraction of the leaf capacity on the y-axis. The worst-case space utilization for the classical B$^+$-tree (and consequently, all other baselines) occurs when ingesting fully sorted data, as every node is only half full due to right-deep insertions. QuIT alleviates this memory overhead by variable splitting guided by *IKR*. With near-sorted data ($1\% \leq K \leq 10\%$), the average leaf occupancy of the classical B$^+$-tree is between 51-54%, while QuIT offers an average leaf occupancy between 62-74%. The variable split strategy in QuIT allows for tightly packing the leaf nodes that receive ordered inserts through the $po\ell e$. Further, redistribution also increases leaf occupancy when $po\ell e\_prev$ is less than half full, utilizing previously unused space. For less sorted or fully scrambled data, QuIT is able to match the leaf occupancy of the classical B$^+$-tree.

QuIT Does Not Incur Any Overhead For Point Lookups. QuIT performs exactly as many accesses during point lookups as a classical B$^+$-tree, as the lookup algorithms are identical. Hence, QuIT does not incur any overhead for point lookups. From Fig. 10b, we observe that point lookups in QuIT can in fact, be slightly faster ($\approx 2\%$ on average) as the tree is smaller due to better space utilization. This results in the overall size of the index nodes in QuIT being smaller than the system's cache, effectively leading to a marginal performance improvements.

QuIT is Faster for Range Lookups. Range lookups are faster in QuIT, as shown in Fig. 10c, when ingesting near-sorted data. We observe that range queries access up to 2× fewer leaf nodes ($\approx 1.3\times$ on average) compared to the B$^+$-tree when the ingested data has high sortedness ($K \leq 10\%$). This benefit gradually declines as sortedness decreases since space utilization in B$^+$-tree also improves. However, even with low data sortedness ($K \geq 25\%$), range lookups in QuIT access 1.15× fewer leaf nodes than the B$^+$-tree, directly correlating to improved performance. Overall, QuIT benefits from the variable split and redistribute strategies that tightly pack the leaf nodes, whereas, the B$^+$-tree experiences its worst-case space amplification when ingesting data with a high degree of sortedness.

5.2 Sensitivity Analysis

We now vary several aspects of the workload and experiment with different data sortedness levels, different data sizes, and different sequences of insertions that stress test fast-path ingestion.

5.2.1 Varying Data Sortedness. First, we compare the performance of QuIT against the $\ell i\ell$-B$^+$-tree as we vary data sortedness. Here, we show the fraction of fast-inserts performed during ingestion in the $\ell i\ell$-B$^+$-tree and QuIT in Fig. 11a-11b, while we indicate the space utilization in both the indexes by measuring the average leaf occupancy in Fig. 11c-11d.

Fast-insertions Decrease as K Increases. We observe in Fig. 11a that the fraction of fast-insertions performed by both the $\ell i\ell$-B$^+$-tree and QuIT decreases as we increase the unordered entries (K) in the workload. While both the indexes are comparable in their ability to perform fast-insertions for highly sorted data, the QuIT's benefits are more pronounced when the ingested data becomes less sorted. For less-sorted data ($K=25\%$) and fairly scrambled data ($K=50\%$), QuIT performs, respectively, 15% and 20% more fast-inserts than the $\ell i\ell$-B$^+$-tree. The latter pays an additional penalty (with a top-insert) for moving the fast-path node back to the ordered position for every out-of-order insertion. QuIT, with a robust $po\ell e$ as the fast-path, avoids these naïve movements.

Fast-insertions are Not Affected by Varying L. We also observe in Fig. 11a and 11b that the fraction of fast-inserts performed is unaffected by the increase in L in the ingested workload. This holds for both $\ell i\ell$-B$^+$-tree and QuIT because the use of fast-path optimization depends on the inserted entry being out-of-order with respect to the range of the target leaf page (for both $\ell i\ell$ and $po\ell e$). Even with small values of L (e.g., $L=1\%$) and uniform sortedness distribution [37], the median displacement of out-of-order entries is $\approx 0.5\%$ of the total data size. This is much larger than the range

Figure 11: Comparing $\ell i\ell$-B$^+$-tree and QuIT when varying data sortedness in the ingestion workload: (a) $\ell i\ell$ performs best with high data sortedness; (b) QuIT maximizes fast inserts when compared to $\ell i\ell$; (c) B$^+$-tree (and $\ell i\ell$-B$^+$-tree) suffers worst-case space amplification for high data sortedness; (d) QuIT outperforms B$^+$-tree's space utilization for any degree of data sortedness.

of entries in a leaf page. Only when L is extremely small are the out-of-order entries likely to fit within the leaf-page's range and can utilize the fast-path optimization.

Space Utilization Increases as K Increases. When comparing the space utilization of the $\ell i\ell$-B$^+$-tree and QuIT in Fig. 11c and 11d, we observe that the average leaf occupancy in $\ell i\ell$-B$^+$-tree increases as the ingestion workload becomes less sorted, while an inverse trend is observed in QuIT. Since $\ell i\ell$ makes no effort to improve the index's memory footprint, the space utilization in $\ell i\ell$-B$^+$-tree is identical to the classical B$^+$-tree. $\ell i\ell$-B$^+$-tree's worst-case space amplification occurs for highly sorted data as almost 50% of the space reserved for future insertions after every split is left largely unused. However, as the data becomes less sorted, this reserved space is useful to accommodate the out-of-order insertions.

Meanwhile, QuIT achieves optimal space utilization for fully sorted data, while outperforming $\ell i\ell$-B$^+$-tree (and in turn, the classical B$^+$-tree) by up to 24% for near-sorted data, due to the variable split and redistribute strategies employed when splitting the fast-path node (i.e., $po\ell e$). The average leaf occupancy in QuIT decreases with lower data sortedness, as the fast-path utilization reduces and top-insertions are more pronounced. This results in QuIT's space utilization being comparable to a B$^+$-tree (and $\ell i\ell$-B$^+$-tree) when the fraction of out-of-order entries in the ingested data is very high. Note that we can also tune QuIT to avoid being 100% full for the fully-sorted data if we anticipate out-of-order entries in the future and we want to avoid propagating splits.

5.2.2 Sensitivity Analysis on Data Size. Next, we analyze the scalability of QuIT by increasing the number of entries ingested into the index from 50M to 4B (scaling the data size from 40MB to 32GB). We pick candidate data collections reflecting three degrees of data sortedness - (i) fully sorted data (K=0%), (ii) nearly-sorted data (K=L=5%) that has a sizeable fraction of out-of-order entries, and (iii) less sorted data (K=L=25%) that has significantly many out-of-order entries. Table 3 summarizes our results when compared to the classical B$^+$-tree. We observe that QuIT's ability to utilize its fast-path and exploit sortedness in the ingested workload remains unaffected as the data size grows. The observed fraction of fast-inserts is 100% for fully sorted data, $\approx$ 95% for nearly-sorted

Sortedness	Metric	Data Size (GB)					
		0.4	2	4	8	16	32
fully sorted	Speedup	3.13×	3.19×	3.23×	3.25×	3.27×	3.31×
	% fast-inserts	100%	100%	100%	100%	100%	100%
nearly sorted	Speedup	2.43×	2.52×	2.56×	2.57×	2.61×	2.77×
	% fast-inserts	95.2%	95.2%	95.2%	95.2%	95.2%	95.2%
less sorted	Speedup	1.31×	1.32×	1.33×	1.34×	1.35×	1.35×
	% fast-inserts	74.6%	74.6%	76.4%	74.6%	74.6%	74.6%

Table 3: QuIT scales with data size.

data, and $\approx$ 75% for less sorted data. This aligns with the expected optimal performance for a sortedness-aware index from Fig. 5b. The speedup during ingestion, however, is slightly amplified when scaling the data size and the number of levels in the tree grows. Scaling data size results in longer tree traversals and, therefore, an increase in indexing cost (top-insert) and maintenance. QuIT continues to employ fast-path insertions enabled by $po\ell e$, offering an amortized indexing cost. Overall, QuIT is at least 2× faster than B$^+$-tree when ingesting near-sorted data.

5.2.3 Stress Testing the Fast-Path. We now explore the performance of the fast-path optimizations in tail-B$^+$-tree, $\ell i\ell$-B$^+$-tree, and $po\ell e$-B$^+$-tree (i.e., QuIT without variable split, redistribute, and reset strategies), along with the full design of QuIT for workloads that alternate between near-sorted and fully scrambled. We anticipate that such a workload will be harder to predict and use it to stress-test all fast-path optimizations and *IKR*. We ingest into each index 25M entries divided into 5 segments (of 5M entries), alternating between near-sorted data (K=10%) or scrambled data (K=100%), while we default L to 100%. We visualize this workload using the position (x-axis) and the value (y-axis) of the keys inserted in Fig. 12a.

Fig. 12b shows the number of entries ingested to the index using the fast path (on y-axis) at different snapshots on the x-axis (i.e., one for each segment). Note that a flat line in any segment implies that the index performed only top-inserts. We observe that the tail-B$^+$-tree very quickly reaches the stale state and fails to perform fast-path insertion as soon as the data becomes even slightly unordered. The $\ell i\ell$-B$^+$-tree continues to perform fast-path insertion when the data is nearly sorted, while only performing
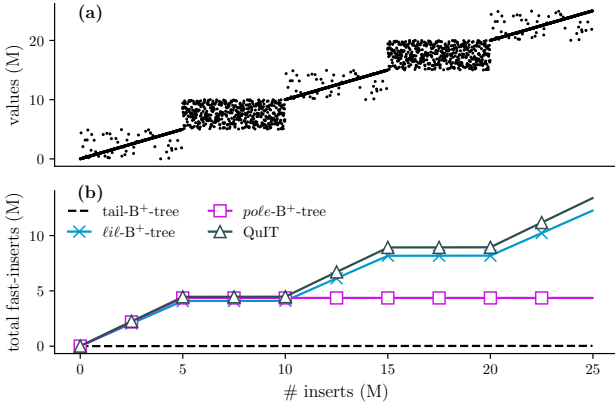
Figure 12: Stress testing the Fast-path optimizations under workloads with varying densities of data sortedness: (a) a workload that varies sortedness density by alternating between near-sorted and scrambled data in different data segments; (b) performance of tail, $\ell i\ell$, and $po\ell e$ optimizations in B$^+$-trees when comparing fast-insertions in QuIT.

top-inserts for the scrambled data segments. Meanwhile, the $po\ell e$-B$^+$-tree marginally outperforms the $\ell i\ell$-B$^+$-tree in the first near-sorted data segment, and subsequently, only performs top-inserts for the remaining workload. This is because $po\ell e$ is trapped in a stale state after the first scrambled segment of data (between keys 5M and 10M). QuIT addresses the limitations of $po\ell e$ and recovers from this stale state with the help of its reset strategy (described in §4.3), allowing continued use of the fast-path when inserting near-sorted data. A robust fast-path ingestion strategy (like $po\ell e$) coupled with the reset strategy helps QuIT outperform $\ell i\ell$-B$^+$-tree by performing $\approx$ 11% more fast-path insertions.

5.3 QuIT under Concurrent Execution

We now benchmark QuIT and compare it with the classical B$^+$-tree as we have more threads concurrently using the two indexes. We experiment with three levels of sortedness as defined already and report the raw throughput in terms of operations per second.

QuIT Scales Better Than B$^+$-tree for Inserts. Ingesting near-sorted data using multiple threads is expected to have high contention because most insertions target the same leaf. As a result, we expect that having multiple threads will hurt performance for both B$^+$-tree and QuIT. However, we anticipate that the lightweight insertion of QuIT's fast path would lead to a shorter critical section. Fig. 13a corroborates our expectations, showing that despite both indexes having contention with a higher number of threads, QuIT offer 1.5-2$\times$ higher throughput than B$^+$-tree in most cases. We also note that QuIT's benefit increases with higher concurrency.

QuIT Scales Better Than B$^+$-tree for Reads. As expected, read queries in both trees behave very similarly since QuIT's read path is essentially the same as with the classical B$^+$-tree. Fig. 13b shows that both trees scale almost perfectly until 8 threads, with the scaling slowing down for 16 threads.

5.4 Comparing with SWARE

We now compare QuIT vs. SA-B$^+$-tree, the state-of-the-art sortedness-aware index that employs the SWARE design [36]. We ingest 500M


Figure 13: QuIT scales better during concurrent execution.

entries (4GB) into both indexes and measure (i) the average latency per insertion; and (ii) the average point lookup latency, when varying the fraction of out-of-order entries in the data collection (defaulting L=100%). Note that we utilize a more optimized B$^+$-tree (also used in QuIT) than the one provided with the SWARE codebase [2] as the underlying tree index in the SA-B$^+$-tree. SWARE originally packs a B$^\epsilon$-tree [6] that also functions as a B$^+$-tree when appropriately tuned. This, however, introduces orthogonal complexities and overheads that we avoid in our B$^+$-tree prototype.

QuIT Outperforms SWARE During Ingestion. Fig. 14a shows that QuIT offers significantly better latency per ingestion than the SA-B$^+$-tree for any data sortedness. Even when ingesting fully sorted data, QuIT offers a 16% improvement over the SA-B$^+$-tree, as every insert can be directly appended to the fast-path node ($po\ell e$). Despite SA-B$^+$-tree performing opportunistic bulk loading on-the-fly, it still pays a cost to index the data in the buffer through two levels of Bloom filters (the global and per-page Bloom filters) and the Zonemaps, in addition to costs associated with data movement in the buffer after every flush operation. Likewise, for near-sorted data ($K\leq$10%), QuIT is at least 1.5$\times$ (and 1.86$\times$ on average) better than SA-B$^+$-tree during ingestion. The additional costs incurred by the SA-B$^+$-tree to update the necessary metadata (i.e., identifying non-overlapping zones, Zonemaps, and Bloom filters) for facilitating opportunistic bulk loading far exceed the costs associated with maintaining QuIT's metadata, primarily due to the latter's lightweight index design. Thus, QuIT's outperforms SA-B$^+$-tree during near-sorted data ingestion. Both indexes perform top-inserts for those entries that cannot utilize the fast-path optimization (opportunistic bulk loading in the case of SA-B$^+$-tree). As data sortedness decreases, top-inserts are more pronounced, thus, the average latency for an insertion increases for both indexes. QuIT and SA-B$^+$-tree are have comparable performance for less sorted ($K\geq$25%) or scrambled data.

QuIT Offers Faster Lookups than SWARE. Fig. 14b compares the lookup performance of SA-B$^+$-tree and QuIT with varying data sortedness. Here, we perform 5M point lookups (generated uniformly and randomly) on the ingested keys for both indexes. QuIT does not incur any read overhead due to its lightweight design, while the SA-B$^+$-tree incurs an additional cost of scanning the buffer for the target key. While SA-B$^+$-tree employs additional data structures like Bloom filters and Zonemaps, as well as techniques like query-driven partial sorting to reduce this cost, it still pays an increased cost compared to QuIT, which simply performs a B$^+$-tree lookup. Overall, QuIT outperforms the SA-B$^+$-tree by up to 26%. Only when querying the index after ingesting fully sorted

Aneesh Raman, Konstantinos Karatsenidis, Shaolin Xie, Matthaios Olma, Subhadeep Sarkar, and Manos Athanassoulis


Figure 14: Comparing the SA-B$^+$-tree and QuIT: (a) A complex design allows SA-B$^+$-tree to opportunistically bulk load near-sorted data while QuIT maximizes fast insertions through index appends; (b) SA-B$^+$-tree incurs a read-overhead while QuIT is marginally faster than the B$^+$-tree.

data, the SA-B$^+$-tree offers an $\approx$ 8% improvement over QuIT as the queries targeting keys still in the SWARE buffer are answered faster, because, for fully sorted data, entries in the buffer can make the most efficient use of Zonemaps and interpolation search [42].

5.5 Indexing Real-World Data

Real-world data frequently exhibit a degree of near-sortedness that may be unknown or difficult to quantify, where the K-L metric may not be a natural descriptor. For example, Fig. 15 shows the `closing_price` data from two stock price instruments: (i) NIFTY, that is the equity benchmark index of the National Stock Exchange, India; and (ii) SPXUSD, that is the S&P 500 index. We obtain the intra-day stock price data at one-minute timeframes (sources listed in footnote[3] and [4]). Both datasets contain $\approx$ 1.4M entries and 2.2M entries, respectively. We see an overall upward trend that intuitively implies near-sortedness.

QuIT Offers Best Performance for Real-World Data. Fig. 15c shows the speedup offered when ingesting the stock price data into the tail-B$^+$-tree, SA-B$^+$-tree, $\ell i\ell$-B$^+$-tree, and QuIT, normalized vs. the baseline B$^+$-tree. QuIT offers a $\approx$ 30% improvement on average during ingestion when compared to the tail-B$^+$-tree. In fact, our approach offers the maximum speedup among all baselines, even outperforming the state-of-the-art sortedness-aware index by $\approx$ 8% and 5% for the NIFTY and SPXUSD instruments, due to its lightweight design. Overall, tree indexes like B$^+$-tree-based designs benefit from sortedness-aware ingestion optimizations even when intrinsic data sortedness is hard to quantify or predict, as shown by the ingestion speedup offered by both SWARE and QuIT.

6 RELATED WORK

There is a plethora of B$^+$-tree variants optimizing for data ingestion. We recognize two design patterns aiming to reduce the cost of incremental data insertion: (i) approaches that optimize tree traversal and insertion operations by taking maximum advantage of modern hardware, and (ii) approaches that re-design the internal structure of the tree to amortize the cost of insertions across the workload. Specifically, the CSB-tree [39] resizes nodes to make all operations cache-conscious and minimize cache misses. The PLI-tree [41], T-tree [26], YATS-tree [23], Partitioned B$^+$-trees [17] and B$^\epsilon$-trees [6] adapt the data layout in nodes to their corresponding use case. For

[3]https://github.com/aeron7/nifty-banknifty-intraday-data
[4]https://github.com/FutureSharks/financial-data

Figure 15: Real-world data can often carry implicit sorted-ness, for example, in closing prices of stock instruments like: (a) NIFTY, and (b) SPXUSD. (c) QuIT offers the best ingestion performance for such near-sorted data.

instance, B$^\epsilon$-trees trade-off fan-out for per-node buffers that can batch insertions. This allows the B$^\epsilon$-tree to amortize the ingestion cost, while the buffers are gradually flushed down the index. Finally, Bw-trees [27] are log-structured and take a latch-free approach using a delta update scheme with dynamically sized pages. The nodes of the Bw-tree are only logical and do not occupy fixed physical locations on main memory or storage. Bw-trees design eliminates thread blocking and is optimized for modern hardware.

These B$^+$-tree designs improve ingestion performance, however, they are unaware of prospective gains opportunities from taking advantage implicit sortedness of the incoming data. The SWARE indexing paradigm takes advantage of sortedness by using a combination of in-memory buffering and bulk-loading to optimize index ingestion [38]. However, its gains require a complex design that utilizes additional resources, adversely affecting the lookup performance. Meanwhile, QuIT offers sortedness-awareness index ingestion through a lightweight design and minimal metadata footprint.

Applicability to Data Streaming and Time Series. Time series indexing assumes that data ingestion follows an expected increasing order [25, 33, 43–45]. Streaming systems often use a buffer to capture the arrival skew within time-based windows [40] that allow for effective data series comparisons [10, 16]. QuIT can eliminate the need for the additional buffer as in-order data will always be fast-inserted, leaving the arrival data skew to be captured by the fraction of top-inserts performed, as we demonstrate in §5.2.3.

7 CONCLUSION

Commercial data systems employ fast-path optimization techniques to amortize the cost of index construction during data ingestion. For B$^+$-trees fast-path ingestion helps avoid tree traversals, inserting entries directly to the tail leaf if the inserted data is fully sorted. However, this tail-leaf fast path becomes stale when the data is not fully sorted. We address this by proposing two new fast-path ingestion strategies for B$^+$-trees – $\ell i\ell$ and $po\ell e$ – that target near-sorted data. Further, we present QuIT, a lightweight index design that reduces the indexing cost proportionally to the sortedness of the indexed data. In addition, QuIT improves the memory footprint of the index, while also offering better lookup performance. Overall, QuIT outperforms the tail fast-path (and SWARE) by up to 2.5× (2×) when ingesting near-sorted data while offering on average 20% better space utilization when compared to the B$^+$-tree. The reduced memory footprint helps QuIT access up to 2× fewer nodes during range lookups, while it does not incur overhead for point lookups.

REFERENCES

[1] Daniar Achakeev and Bernhard Seeger. 2013. Efficient Bulk Updates on Multi-version B-trees. *Proceedings of the VLDB Endowment* 6, 14 (2013), 1834–1845.

[2] Manos Athanassoulis and Anastasia Ailamaki. 2014. BF-Tree: Approximate Tree Indexing. *Proceedings of the VLDB Endowment* 7, 14 (2014), 1881–1892.

[3] Manos Athanassoulis, Michael S. Kester, Lukas M. Maas, Radu Stoica, Stratos Idreos, Anastasia Ailamaki, and Mark Callaghan. 2016. Designing Access Methods: The RUM Conjecture. In *Proceedings of the International Conference on Extending Database Technology (EDBT)*. 461–466.

[4] Rudolf Bayer and Karl Unterauer. 1977. Prefix B-trees. *ACM Transactions on Database Systems (TODS)* 2, 1 (1977), 11–26.

[5] Sagi Ben-Moshe, Yaron Kanza, Eldar Fischer, Arie Matsliah, Mani Fischer, and Carl Staelin. 2011. Detecting and Exploiting Near-Sortedness for Efficient Relational Query Evaluation. In *Proceedings of the International Conference on Database Theory (ICDT)*. 256–267.

[6] Michael A. Bender, Martin Farach-Colton, William Jannen, Rob Johnson, Bradley C. Kuszmaul, Donald E. Porter, Jun Yuan, and Yang Zhan. 2015. An Introduction to Bε-trees and Write-Optimization. *White Paper* (2015).

[7] Philip A. Bernstein, Vassos Hadzilacos, and Nathan Goodman. 1987. *Concurrency Control and Recovery in Database Systems*. Addison-Wesley.

[8] Timo Bingmann. 2007. STX B+ Tree. *https://github.com/bingmann/stx-btree* (2007).

[9] Burton H Bloom. 1970. Space/Time Trade-offs in Hash Coding with Allowable Errors. *Commun. ACM* 13, 7 (1970), 422–426.

[10] Paris Carbone, Marios Fragkoulis, Vasiliki Kalavri, and Asterios Katsifodimos. 2020. Beyond Analytics: The Evolution of Stream Processing Systems. In *Proceedings of the ACM SIGMOD International Conference on Management of Data*. 2651–2658.

[11] Svante Carlsson and Jingsen Chen. 1992. On Partitions and Presortedness of Sequences. In *Acta Informatica*, Vol. 29. 267–280.

[12] Douglas Comer. 1979. The Ubiquitous B-Tree. *Comput. Surveys* 11, 2 (1979), 121–137.

[13] CouchDB. [n. d.]. Online reference. *http://couchdb.apache.org/* ([n. d.]).

[14] Frederik M. Dekking, Cornelis Kraaikamp, Hendrik P. Lopuhaä, and Ludolf E. Meester. 2005. *A Modern Introduction to Probability and Statistics*. Springer London. 488 pages.

[15] Jochen Van den Bercken and Bernhard Seeger. 2001. An Evaluation of Generic Bulk Loading Techniques. In *Proceedings of the International Conference on Very Large Data Bases (VLDB)*. 461–470.

[16] Marios Fragkoulis, Paris Carbone, Vasiliki Kalavri, and Asterios Katsifodimos. 2020. A Survey on the Evolution of Stream Processing Systems.

[17] Goetz Graefe. 2003. Sorting And Indexing With Partitioned B-Trees. In *Proceedings of the Biennial Conference on Innovative Data Systems Research (CIDR)*.

[18] Goetz Graefe. 2010. A survey of B-tree locking techniques. *ACM Transactions on Database Systems (TODS)* 35, 3 (2010).

[19] Goetz Graefe. 2011. Modern B-Tree Techniques. *Foundations and Trends in Databases* 3, 4 (2011), 203–402.

[20] Stratos Idreos and Mark Callaghan. 2020. Key-Value Storage Engines. In *Proceedings of the ACM SIGMOD International Conference on Management of Data*. 2667–2672.

[21] Stratos Idreos, Martin L. Kersten, and Stefan Manegold. 2007. Database Cracking. In *Proceedings of the Biennial Conference on Innovative Data Systems Research (CIDR)*.

[22] Stratos Idreos, Stefan Manegold, Harumi Kuno, and Goetz Graefe. 2011. Merging What's Cracked, Cracking What's Merged: Adaptive Indexing in Main-Memory Column-Stores. *Proceedings of the VLDB Endowment* 4, 9 (2011), 586–597.

[23] Chris Jermaine, Anindya Datta, and Edward Omiecinski. 1999. A Novel Index Supporting High Volume Data Warehouse Insertion. In *Proceedings of the International Conference on Very Large Data Bases (VLDB)*. 235–246.

[24] Donald E. Knuth. 1997. *The art of computer programming, Volume I: Fundamental Algorithms (3rd Edition)*. Addison-Wesley.

[25] Haridimos Kondylakis, Niv Dayan, Kostas Zoumpatianos, and Themis Palpanas. 2018. Coconut: A Scalable Bottom-Up Approach for Building Data Series Indexes. *Proceedings of the VLDB Endowment* 11, 6 (2018), 677–690.

[26] Tobin J. Lehman and Michael J. Carey. 1986. A Study of Index Structures for Main Memory Database Management Systems. In *Proceedings of the International Conference on Very Large Data Bases (VLDB)*. 294–303.

[27] Justin J. Levandoski, David B. Lomet, and Sudipta Sengupta. 2013. The Bw-Tree: A B-tree for New Hardware Platforms. In *Proceedings of the IEEE International Conference on Data Engineering (ICDE)*. 302–313.

[28] Heikki Mannila. 1985. Measures of Presortedness and Optimal Sorting Algorithms. *IEEE Transactions on Computers (TC)* 34, 4 (1985), 318–325.

[29] Guido Moerkotte. 1998. Small Materialized Aggregates: A Light Weight Index Structure for Data Warehousing. In *Proceedings of the International Conference on Very Large Data Bases (VLDB)*. 476–487.

[30] MongoDB. 2023. Online reference. *http://www.mongodb.com/* (2023).

[31] MySQL. 2023. MySQL. *https://www.mysql.com/* (2023).

[32] Oracle. 2018. Introducing Oracle Database 18c. *White Paper* (2018).

[33] Themis Palpanas. 2015. Data Series Management: The Road to Big Sequence Analytics. *ACM SIGMOD Record* 44, 2 (2015), 47–52.

[34] PostgreSQL. 2023. PostgreSQL: The World's Most Advanced Open Source Relational Database. *https://www.postgresql.org* (2023).

[35] Raghu Ramakrishnan and Johannes Gehrke. 2002. *Database Management Systems*. McGraw-Hill Higher Education, 3rd edition.

[36] Aneesh Raman, Konstantinos Karatsenidis, Subhadeep Sarkar, Matthaios Olma, and Manos Athanassoulis. 2022. BoDS: A Benchmark on Data Sortedness. In *Proceedings of the TPC Technology Conference on Performance Evaluation \& Benchmarking (TPCTC)*.

[37] Aneesh Raman, Subhadeep Sarkar, Matthaios Olma, and Manos Athanassoulis. 2022. OSM-tree: A Sortedness-Aware Index. *CoRR* abs/2202.0 (2022).

[38] Aneesh Raman, Subhadeep Sarkar, Matthaios Olma, and Manos Athanassoulis. 2023. Indexing for Near-Sorted Data. In *Proceedings of the IEEE International Conference on Data Engineering (ICDE)*.

[39] Jun Rao and Kenneth A. Ross. 2000. Making B+-trees Cache Conscious in Main Memory. In *Proceedings of the ACM SIGMOD International Conference on Management of Data*. 475–486.

[40] Utkarsh Srivastava and Jennifer Widom. 2004. Flexible Time Management in Data Stream Systems. In *Proceedings of the ACM SIGACT-SIGMOD-SIGART Symposium on Principles of Database Systems (PODS)*. 263–274.

[41] Kristian Torp, Leo Mark, and Christian S Jensen. 1998. Efficient Differential Timeslice Computation. *IEEE Trans. Knowl. Data Eng.* 10, 4 (1998), 599–611.

[42] Peter Van Sandt, Yannis Chronis, and Jignesh M. Patel. 2019. Efficiently Searching In-Memory Sorted Arrays: Revenge of the Interpolation Search?. In *Proceedings of the ACM SIGMOD International Conference on Management of Data*. 36–53.

[43] Kostas Zoumpatianos, Stratos Idreos, and Themis Palpanas. 2014. Indexing for interactive exploration of big data series. In *Proceedings of the ACM SIGMOD International Conference on Management of Data*. 1555–1566.

[44] Kostas Zoumpatianos, Stratos Idreos, and Themis Palpanas. 2016. ADS: the adaptive data series index. *The VLDB Journal* 25, 6 (2016), 843–866.

[45] Kostas Zoumpatianos and Themis Palpanas. 2018. Data Series Management: Fulfilling the Need for Big Sequence Analytics. In *Proceedings of the IEEE International Conference on Data Engineering (ICDE)*. 1677–1678.